Maarij Aqeel

Faisalabad, Pakistan maarijaqeel3200@gmail.com +92 3328885574
github.com/Maarij-Aqeel linkedin.com/in/maarijaqeel

SUMMARY

AI Engineer with 2 years of experience architecting production LLM applications and scalable backend systems, specializing in RAG pipelines, fine-tuned language models, and API-driven architectures serving **10K+** daily users. Proven track record of deploying AI solutions that reduce operational costs by **75%+** and accelerate workflows from days to hours.

TECHNICAL SKILLS

Languages:	Python, JavaScript, TypeScript, SQL
AI & MLOps:	RAG Pipelines, **LangChain**, **Ollama**, HuggingFace, Fine-Tuning (**LoRA**/QLoRA), **Vapi**, PyTorch, TensorFlow
Backend & Cloud:	**FastAPI**, Django, PostgreSQL, MongoDB, Redis, Azure, AWS, Docker, CI/CD, Git
Frontend:	**Next.js 16**, **React 19**, Tailwind CSS, RESTful API Design
Automation:	**n8n**, Make, Zapier, Workflow Automation, Agentic Design

EXPERIENCE

Contract AI Engineer — *Maravilla Cleaners, 16x9 Studio*

2024 – Present — *Remote*

- Architected end-to-end Voice + Chat AI agent for Maravilla Cleaners using **Vapi Voice API** and **n8n**, automating appointment scheduling and reducing manual booking overhead by **90%** while capturing **40%+** more after-hours leads.
- Engineered AI automation pipeline for 16x9 Studio using **n8n** and LLM integrations, reducing project delivery time from **5 days to 1.5 days** (**70%** acceleration).
- Deployed multi-tenant private cloud using OpenStack and Ansible, cutting developer setup time from **4 hours to 10 minutes** (**96%** reduction).

Backend & AI Engineer — *Software Productivity Strategists (SPS)*

Sep 2024 – Dec 2024 — *Remote*

- Productionized **FastAPI**-based REST endpoints serving **10K+** daily requests, optimizing PostgreSQL queries and implementing Redis caching to cut API response time from **450ms to 315ms** (**30%** improvement).
- Integrated IBM WatsonX LLM into Greenkeeper app, engineering RAG pipelines that automated **12-step** approval workflows, reducing project approval time from **3 days to 4 hours** (**90%** reduction) for **50+** concurrent projects.

AI Development Engineer — *Nexium*

Jun 2024 – Aug 2024 — *Remote*

- Architected Intervue AI platform using **Next.js**, **Vapi Voice API**, and **n8n** with dual-role authentication and real-time analytics, reducing interview setup time by **30%**.
- Productionized AI blog summarization pipeline using Gemini and **LangChain**, processing **200+** articles weekly and reducing content review cycle by **75%** (**2 hours to 30 minutes**).

PROJECTS

RoleSync AI – ATS Resume Optimizer & Cover Letter Generator — ***Next.js 16**, **React 19**, Google Gemini 2.5 Flash, TipTap, Zustand, LaTeX*

- Architected AI-powered resume optimization platform scoring resumes against JDs using dual-metrics (ATS + Domain Fit), achieving **85%+** user-reported interview callback improvement.
- Engineered real-time keyword injection with DOM-scanning AI and TipTap editor, enabling one-click insertion of metric-driven bullet points with visual pulse animations.
- Implemented dynamic **LaTeX** PDF export pipeline bridging frontend TipTap to TexLive backend, delivering ATS-optimized resumes without external database persistence.
- Built zero-upload cover letter generator with Zustand state inheritance, featuring ultra-strict prompt engineering for sub-**300-word** technical recruiter-persona outputs and **3-tier** Clipboard API fallback.

Semantic Image Search – Multimodal Retrieval Engine — *OpenAI CLIP, Pinecone, COCO Dataset*

- Engineered multimodal search pipeline using OpenAI CLIP for embedding generation and Pinecone for vector indexing, achieving **<100ms** retrieval latency across **50,000+** image dataset.
- Implemented text-to-image, image-to-image, and hybrid search workflows, enabling nuanced discovery of semantically relevant visual content.

GhostShell – AI-Powered Terminal Assistant — *Python, HuggingFace, DeepSeek-Coder-V2-Base, **LoRA***

- Fine-tuned DeepSeek-Coder-V2 model using **LoRA** on **15K+** command-context pairs, improving command prediction accuracy by **27%** over baseline.
- Built CLI autocompletion system leveraging user history and environment variables to provide intelligent command suggestions.

SmartDietPlanner – AI-Powered Nutrition Assistant — *Django, Botpress, HuggingFace, **LoRA**, Ollama, Pinecone*

- Fine-tuned LLaMA 3.2 using **LoRA** on **35,000** recipe-nutrition samples, achieving **27%** improvement in response relevance and nutritional accuracy.
- Deployed RAG pipeline with Pinecone vector database for real-time recipe retrieval, serving queries with **<2s** response time.

EDUCATION

FAST National University of Computer and Emerging Sciences (FAST NUCES) — *Chiniot, Pakistan*

Bachelor of Science in Artificial Intelligence — Expected Graduation: 2027